

Partitioning3:

Re-balancing partitions:

We need to re-balance partitions in the following conditions:

- Add more machines
- Add more disks
- Replace machines
- Skewed + hotspot partitions

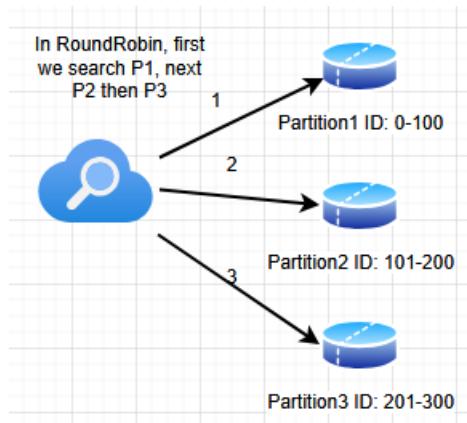
Request Routing:

- Reading mechanisms

Our problem statement is, we have to read from multiple partitions

Approaches are Round Robin,

Round Robin:

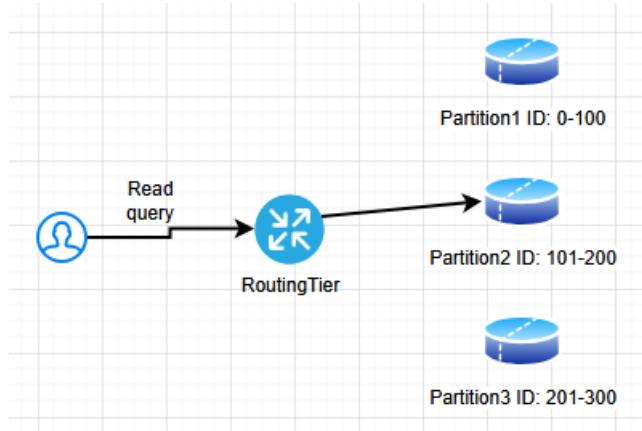


We are checking each and every partition in round-robin manner and once we find the data we stop searching. Time complexity is $O(n)$. It will wait after searching one partition before moving to the next partition.

Demerit is Time complexity is high, latency because we go into all partitions

Routing Tier:

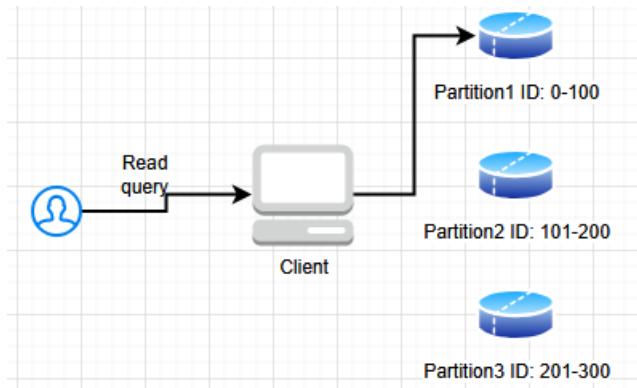
RoutingTier will have all the addresses, it's like a global index. It redirects user read query to the correct partition.



It's better than RoundRobin because Time performance improves. Time complexity O(1)

- RoutingTier determines which partition to look for
- RoutingTier contains logic for routing and maintains metadata for mapping

3. Clients are aware of partitioning and assignment of partitions in nodes:
We remove RoutingTier and client will be managing routing to partitions.



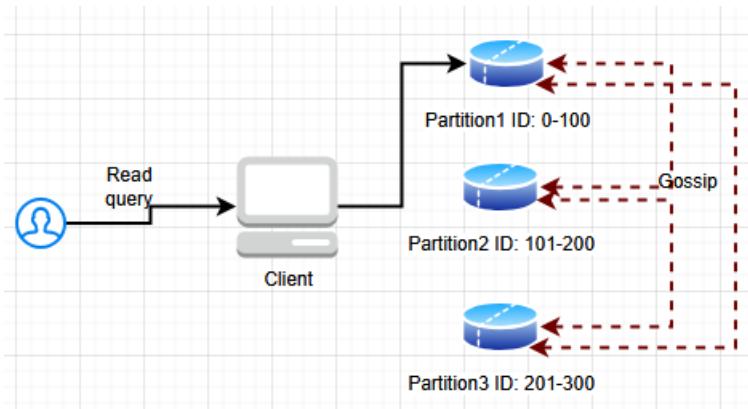
Client is aware that ID: 56 is in Partition1

- No RoutingTier
- All transactions will be handled by our client
- when a new partition arrives, client must be updated about it

Use coordination services: These services keep track of the partition. Examples where coordination services is part of the application are Zookeeper, HBase, Kafka

4. Gossip protocol: Cassandra

First lets say the Read query goes to P1 by default, if it doesn't have the data then nodes/partitions gossip/interact among themselves to locate the data then the query will be redirected to that partition.



Homework:

CouchBase: it is a system from which we can handle data, replication, partitioning
How CouchBase handles routing? Look it up